

Introduction to AWS DynamoDB

Omid Vahdaty, Big Data ninja.



Concepts

- Schema less - other than the primary key attributes, you do not need to define any attributes or data types at table creation time.
- Table == collection/Table from OLTP
- Item == "Row in a table".
- Attributes == "col of a row " in a table. May change from item to item.
- Primary key ==a primary key uniquely identifies each item in the table, so that no two items can have the same key. 2 types
 - **Partitions key** - simple **hash** function that maps **item into specific partition**
 - **Partition + sort key** - A composite primary key
 - Hash function maps item into partition. Also called ***hash attribute***
 - The item will be sorted via sort key. Also called **range attribute**.
- **Secondary Index. Non key index.** Another way to read data. **Faster? :)**
 - Global secondary index – an index with a partition key and sort key that can be different from those on the table.
 - Local secondary index – an index that has the same partition key as the table, but a different sort key.

DynamoDB stream

Optional feature : Each event is represented by a *stream record*. If you enable a stream on a table, DynamoDB Streams will write a stream record whenever one of the following events occurs:

If a new item is added to the table, the stream captures an image of the entire item, including all of its attributes.

If an item is updated, the stream captures the "before" and "after" image of any attributes that were modified in the item.

If an item is deleted from the table, the stream captures an image of the entire item before it was deleted.

Lifetime of 24 hours

Actions

- a. `ListStreams`
- b. `DescribeStream`
- c. `GetShardIterator`
- d. `GetRecords`



DynamoDB API

- Control Plane: create and manage tables:

- CreateTable
- DescribeTable
- ListTables
- UpdateTable
- DeleteTable

- Data Plane : actions on data:

- create, read, update, and delete



DynamoDB API

- Data Plane : actions on data:
 - Create
 - putItem - only 1 item, must specify primary key
 - batchWriteItem - write upto 25 item in transaction
 - Read
 - getItem - get only 1 item
 - batchGetItems - get upto 100 items.
 - Query
 - Scan
 - Update - updateItem, 1 item, specify primary key, atomic counters to save writes? ;)
 - Delete, deleteItem upto 1 item, BatchDeleteItems. (upto 25)

DynamoDB data types

Scalar Types – A scalar type can represent **exactly one value**. The scalar types are number, string, binary, Boolean, and null.

Document Types – A document type can represent a **complex structure with nested attributes**—such as you would find in a JSON document. The document types are list and map.

Set Types – A set type can represent **multiple scalar values**. The set types are string set, number set, and binary set.



Documents?

- complex data structures up to 32 levels deep.
- size limit (400 KB).
- List:
 - A list type attribute can store an ordered collection of values
 - E.g `FavoriteThings: ["Cookies", "Coffee", 3.14159]`
- Map - A map type attribute can store an unordered collection of name-value pairs, ideal for json.

- ```
{
 Day: "Monday",
 UnreadEmails: 42,
 ItemsOnMyDesk: [
 "Coffee Cup",
 "Telephone",
 {
 Pens: { Quantity : 3},
 Pencils: { Quantity : 2},
 Erasers: { Quantity : 1}
 }
]
}
```

# Sets?

- item size limit (400 KB).
- DynamoDB does not support empty sets.
- **Example (String Set, Number Set, and Binary Set)**
- `["Black", "Green", "Red"]`  
`[42.2, -19, 7.5, 3.14]`  
`["U3Vubnk=", "UmFpbmk=", "U25vd3k="]`



# Provision IO

Read carefully:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>

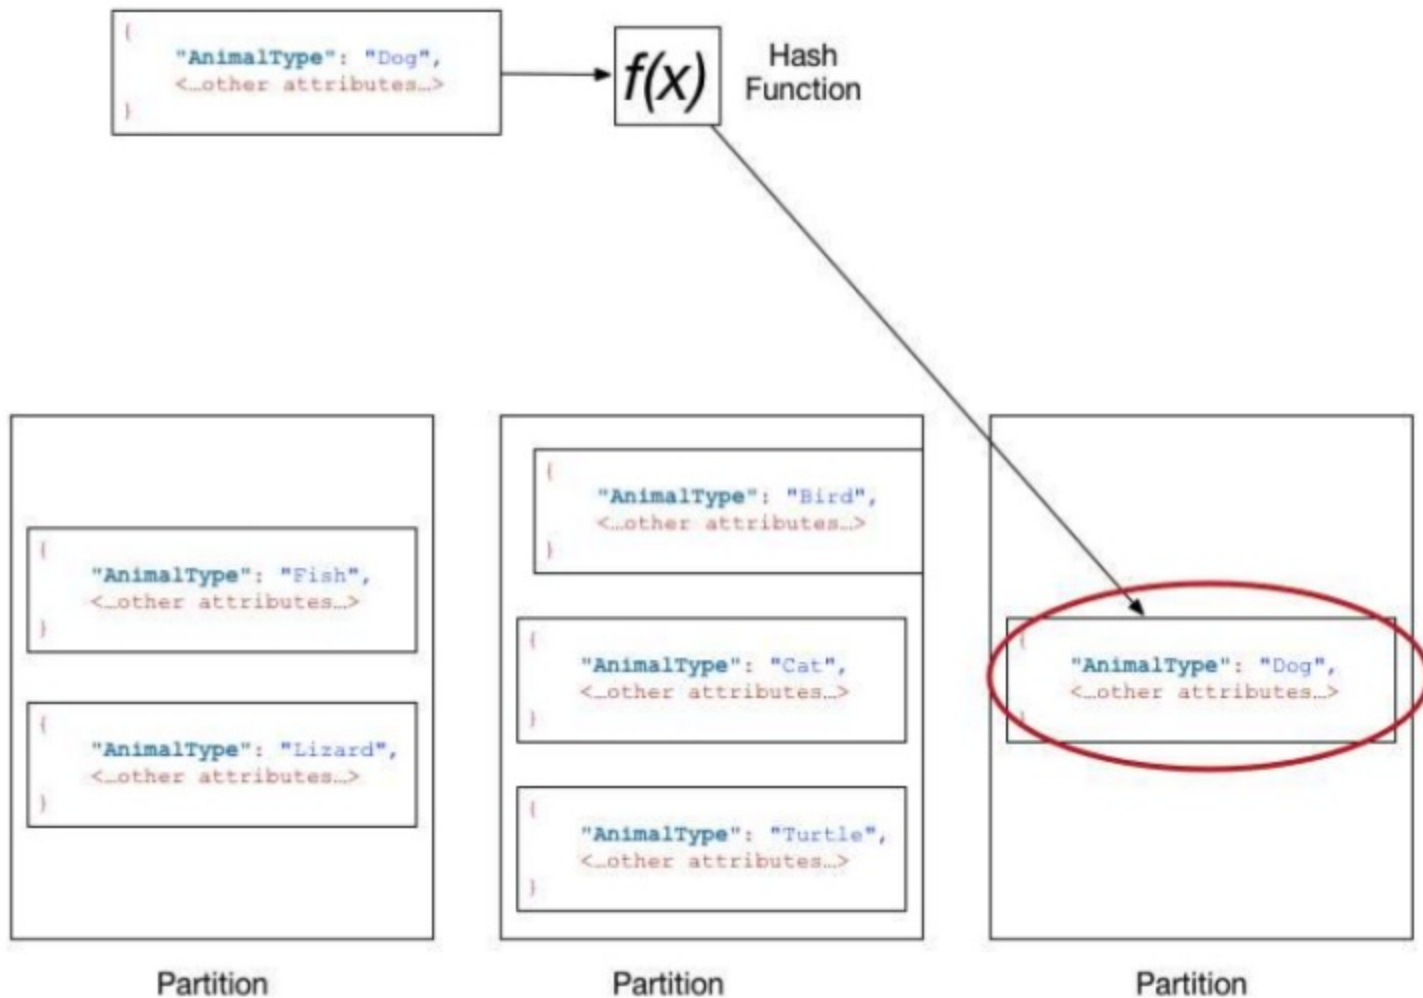


# Partitions?

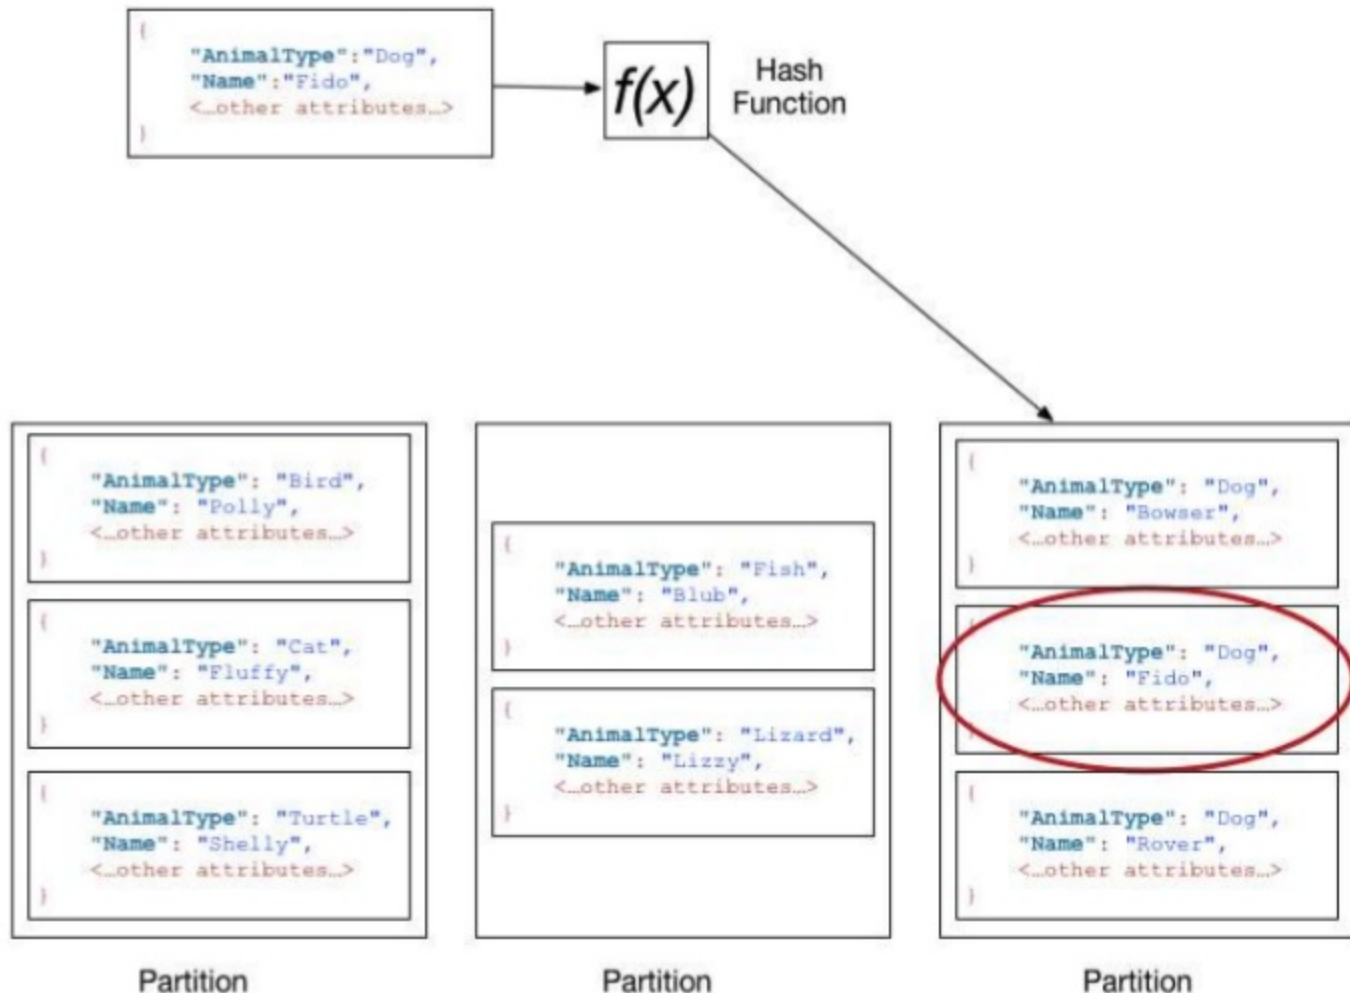
- DynamoDB stores data in partitions.
- **A *partition* is an allocation of storage for a table**
- backed by solid-state drives (SSDs)
- **Multi AZ - automatically** replicated across multiple Availability Zones within an AWS region.
- Partition management is handled automatically by DynamoDB
- Partition key in a table - hash map function to map an item to a partition



# Partitions key



# Partitions key Sort Key



# Getting started

- Download(!) dynamo to your laptop for training:

<http://docs.aws.amazon.com/amazondynamodb/latest/gettingstartedguide/GettingStarted.Download.html>

- Start server

- `java -Djava.library.path=./DynamoDBLocal_lib -jar DynamoDBLocal.jar -sharedDb -inMemory`
- Port 8000



# Java and DynamoDB

Setup access key , secret key

Setup JAVA and dynamo DB:

<http://docs.aws.amazon.com/sdk-for-java/v1/developer-guide/setup-install.html>

Easy to get started with Eclipse, didnt find anything supported for IntelliJ

<https://www.eclipse.org/downloads>

Setup AWS tool kit

<http://docs.aws.amazon.com/toolkit-for-eclipse/v1/user-guide/setup-install.html>

Setup Java SDK:

<http://docs.aws.amazon.com/sdk-for-java/v1/developer-guide/setup-install.html>



# Java and DynamoDB

Using the SDK:

add the full path to the **lib** and **third-party** directories to the dependencies in your build file, and add them to your java **CLASSPATH** to run your code.

Working with local (onsite) installation of dynamodb:

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/DynamoDBLocal.html>

```
o java -Djava.library.path=./DynamoDBLocal_lib -jar DynamoDBLocal.jar -sharedDb
```

Command line options for local DynamoDB (in memory, path to save files etc):

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/DynamoDBLocal.html>

Specify local end point example:

<https://github.com/aws/aws-sdk-java/issues/330>



# Sample Code

```
final String USAGE = "\n" + "Usage:\n" + " CreateTable <table>\n\n" + "Where:\n" + " table - the table to create.\n\n" + "Example:\n" + "CreateTable Hello\n\n";

if (args.length < 1) {

 System.out.println(USAGE);


 System.exit(1);

} /* Read the name from command args */

String table_name = args[0];

System.out.format(

 "Creating table \"%s\" with a simple primary key: \"Name\".\n", table_name);
```





# Sample Code

```
CreateTableRequest request = new CreateTableRequest()

 .withAttributeDefinitions(new AttributeDefinition("Name", ScalarAttributeType.S))

 .withKeySchema(new KeySchemaElement("Name", KeyType.HASH))

 .withProvisionedThroughput(new ProvisionedThroughput(new Long(1), new Long(1)))

 .withTableName(table_name);
```

```
final AmazonDynamoDBClient dynamoDB = new AmazonDynamoDBClient();
```

```
dynamoDB.setEndpoint("http://127.0.0.1:8000");
```

```
dynamoDB.setSignerRegionOverride("eu-west-1");
```



# Sample Code

```
try {

 CreateTableResult result = dynamoDB.createTable(request);

} catch (AmazonServiceException e) {

 System.err.println(e.getMessage());

 System.exit(1);

}

System.out.println("Done!");
```

