

Lessons learned from designing a QA automation for analytics databases (Big Data)

Omid Vahdaty, Big Data ninja.



Agenda

- Introduction to Big Data Testing
- Methodological approach for QA Automation road map
- True Story
- The Solution?
- Challenges of Big Data Testing.
- Performance & Maintenance.



Introduction to Big Data Testing




Scale OUT VS. Scale UP systems




Big Data ?

- Structured (SQL, tabular, strings, numbers)
- Unstructured (logs, pictures, binary, json, blob, video etc)

"80% of business-relevant information originates in unstructured form, primarily text."



Structured Data vs. Unstructured Data



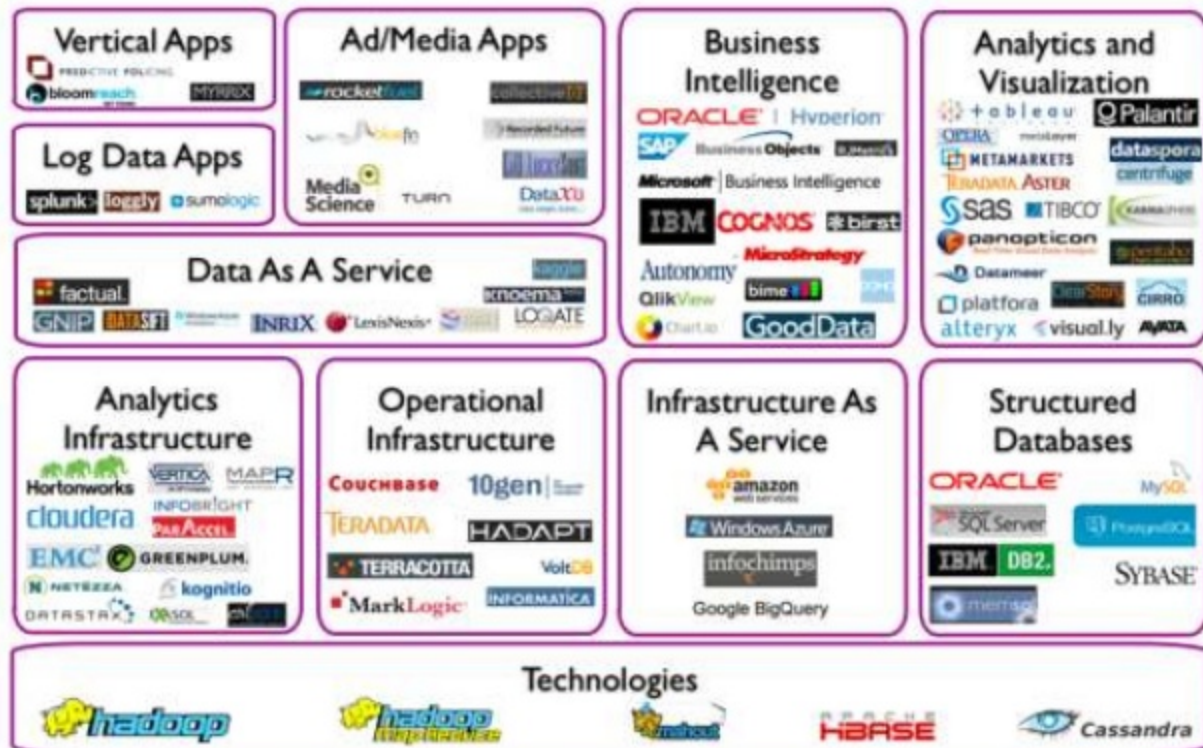
OLAP or OLTP?

- Simply put:
 - **One query** running on a **lot of data** for reporting/analytics
 - **Many queries** running quickly in parallel applications (e.g user login to a website, credentials are stored on a DB)



Type of Big Data products

- Databases
 - OLAP, OLTP
 - Sql, NO SQL
 - In Memory, Disk based.
 - Hadoop Ecosystem, Spark
- Ecosystem tools
 - ETL tools
 - Visualizations tools
- General Application
- Analytics Based products
 - Froud
 - Cyber
 - Finances and etc.



Expectation Matching for today's lecture...

- OLAP Database
- Structured data only, Synthetic data was used in testing.
- (not about QA of analytics based products/solutions).
- (not about hadoop , event streaming cluster).
- (not try to sell anything)
- In a nutshell:
 - How to test an SQL based database?
 - What sort of challenges were met?



How Hard Can it be (aggregation)?

- Select Count(*) from t;
 - Assume
 - 1 Trillion records
 - ad-hoc query (no indexes)
 - Full Scan (no cache)
 - Challenges
 - Time it takes to compute?
 - IO bottleneck? What is IO pattern?
 - CPU bottleneck?
 - Scale up limitations?




How Hard Can it be (join)?

- Select * from t1 join t2 where t1.id = t2.id
 - Assume
 - 1 Trillion records both tables.
 - ad-hoc query (no indexes)
 - Full Scan (no cache)
 - Challenges
 - Time it takes to compute?
 - IO bottleneck? What is IO pattern?
 - CPU bottleneck?
 - Scale out limitations?



3 Pain Points in Big Data DBMS



Big Data ingress challenges

- Parsing of Data type
 - Strings
 - Dates
 - Floats
- Rate of data coming into the system
- ACID: Atomicity, Consistency, Isolation, Durability
- Compression on the fly (non unique values)
- amount of data
- On the fly analytics
- Time constraints (target: x rows per sec/hour)



Big Data egress challenges

- Sort
- Group by
- Reduce
- Join (sortMergeJoin)
- Data distribution
- Compression
- Theoretical Bottlenecks of hardware.



Methodological approach for QA Automation road map



Business constraints?

- Budget?
- Time to market?
- Resources available?
- Skill set available ?



Product requirements?

- What are the product's supported use cases?
- Supported Scale?
- Supported rate of ingress?
- Supported rate on egress?
- Complexity of Cluster?
- High availability?



The Automation: Must have requirements?

- Scale out/up?
- Fault tolerance!
- Reproducibility from scratch!
- Reporting!
- “Views” to avoid duplication for testing?
- Orchestration of same environment per Developer!
- Debugging options
- versioning!!!



The method

- Phase 0: **get ready**
 - Product requirements & business constraints
 - Automation requirements
 - Creating a testing matrix based on insight and product features.
- Phase 1: **Get insights (some baselines)**
 - Test **Ingress** separately, find your **baseline**.
 - Test **Egress** separately, find your **baseline**.
 - Test **Ingress while Egress** in progress., find your **baseline**.
 - Stability and Performance
- Phase 2: Agile: **Design an automation system** that satisfies the requirements
 - **Prioritize** automation features based on your insights from phase 1.
 - **Implement** an automation infrastructure as soon as possible.
 - Update your **testing matrix** as you go (insight will keep coming)
 - **Analyze** the test results daily.
- Phase 3: **Cost reduction**
 - How to reduce compute time/ IO pattern/ network pattern/storage **footprint**
 - **Maintenance** time (build/package/deploy/monitor)
 - **Hardware** costs



**KEEP
CALM
AND USE THE
SCIENTIFIC
METHOD**

True Story

Testing GPU based DBMS

SQL

(Peta Scale)



So why build a DBMS from scratch?

- Most **compiler** are designed for **CPU** → **execution tree** for CPU runtime engine → **new compiler** with **GPU resource** in mind.
- Most **Algorithm** were written for **CPU** → same algorithm logic, **redesign** for **parallelism** philosophy of GPU → **new runtime** with **GPU** resource in mind → **performance increase by order of magnitude**.
- **VISION: fastest DBMS, cost effective, true scalability.**



True story

- Product: Big Data gpu based DBMS system for analytics. [peta scale]
 - Structed data only, designed for OLAP use cases only
- Technical Constraints
 - How to test SQL syntax?(infinite input) **SQL Coverage** from 0% to 20% to 90% to 100%?.
 - What is the expected **impact of big data**?
 - How to **debug performance** issues?
 - Can you **virtualize**? Should you virtualize?
 - **Running time** - how much it takes to run all the tests?
- Company Challenges
 - **Cost** of Hardware? High end 2U Server
 - **Expertise** ? skillset?
 - **Human resources**: Size of QA Automation team?
 - How do you **manage automation**?
 - What is your **MVP**? (Chicken and Egg)
 - **TIME TO MARKET!!!** The product needs to be released.



BASED ON A
TRUE STORY


The baseline concept

- Requirements

- CSV , DDL , query
- Competing DB
- Your DB
- Synthetic data used mostly.
- (real data in peta scale is hard to come by)

- Steps

- Insert same data to same table on both DB's
- Run same query on both DB's
- Compare results on both.
- If equal → test pass.




The Solution: SQL testing framework.

- **Lightweight** python test suite per feature
 - **DDL**
 - **Set of Queries**
 - **Json** for data generation requirements
 - Data generation utilities (**genUtil**)
 - **test results** report - in CSV format.
 - **Expected error** mechanism.
 - **Results compare** mechanism
 - **Command line** arguments for advanced testing/config/tuning/views.
- Wrapper test suite
 - group set of test suites by logic - e.g **Daily**
 - Aggregate **reporting** mechanism
- Scheduling ,reporting,monitoring,deployment, alerting mechanism



Testing concept of SQL syntax

	Simple	Aggregation	Joins
Simple	x	x	x
Aggregation	x	x	x
Joins	x	x	x



Challenges with SQL syntax testing

- (Binary data not supported in the product)
- **Repetition** of queries.
- Different **data type** names.
- Different **data ranges** per data types.
- Different **accuracy** per data types.
- The competition DB that supports big data is expensive...
- **Accuracy** of results. (different DB's return different accuracy)
- **SQL has some extreme cases** (differ per vendor).
- **Datetime** format.
- **Unsupported** features.
- **Duplication** of testing.
- **Very hard to predict which queries are useful** (negative and positive testing).



Challenges with Small Data testing

- Generating Random data
 - **Reproducible** every time --> **Strict data ranges per test** (length, numeric range, format, accuracy)
 - What is the amount of unique values? Different histogram, different bottlenecks:
 - **Unique values challenges:**
 - Per chunk?
 - Per Column
 - **Non unique values challenges:**
 - **String?**
 - Lengths
 - Compressible?
 - **Numeric?**
 - Overflow?
 - Floating point?



Testing matrix: Data Integrity on Big data.

	Data (no null)	Compression	Null Data	Lookup.	Partitions	JDBC/ODBC
1 row						
1M Rows						
10M rows						
100M Rows						
1B rows						
10B rows						
100B rows						
1 Trillion rows						
10 Trillion rows						



Challenges with Big Data testing.

- **Generating** Time of data → genUtil with json for user input
- **Reproducibility** → genUtil again.
- **Disk space** for testing data → peta scale product → peta scale testing data
- **Results compared** mechanism on big data require a big data solution → SMJ
- **Network** bottlenecks on a scale out system...
- **ORDER IS NOT GUARANTEED!**



Insights on the fly: User Defined Test Flow

- A **series of queries in a specific order** - **crushed** the system
- **Solution:** Automation infrastructure on a group of queries: **User defined Test Flow.**
- Good for: pre and post condition of , custom testing, system perspective testing. (partitions)




The solution: some Metrics

1. Extra large **Sanity** Testing per commit (!)
2. **Hourly** testing on latest commit (24 X 7)
 - a. 800 queries
 - b. Zero data
3. **Daily** test:
 - a. 400,000 queries
 - b. Zero data for 90% , 10% testing upto 1B records.
4. **Weekly** testing:
 - a. Big data testing 10B and above.
5. **Monthly** testing:
 - a. Full regression per version.



The solution: Hardware perspective

- 50 x **Desktops**: Optiplex 7040, I7 4 cores, 32GB. 1X 4Tb.
- 10 x **High End Servers**: Dell R730/R720: 20 cores, 128GB. 16X 1TB disks.
Nvidia K40
- DDN **storage** with 200TB+ GPFS
- **Mellanox** FDR switch for the storage
- **1GB switches** for the desktop compute nodes.



Performance and Maintenance



Performance Challenges: app perspective

- **Architecture bottlenecks** [count(*), group by, compression, sort Merge Join]
- What is **IO pattern**?
 - OLTP VS OLAP.
 - Columnar or Row based?
 - How big is your data?
 - READ % vs WRITE %.
 - Sequential? random?
 - Temporary VS permanent
 - Latency VS. throughput.
 - Multi threaded ? or single thread?
 - Power query? Or production?
 - Cold data? Host data?



Performance Challenges: OPS perspective

- Metrics
 - What is Expected Total **RAM required per Query**?
 - OS **RAM cache** , **CPU Cache** hits ?
 - **SWAP** ? yes/no how much?
 - OS metrics - **open files**, stack size, realtime priority?
- Theoretical limits?
 - **Disk type selection** -limitation, expected throughput
 - **Raid** controller, RAID type? Configuration? caching?
 - **File system** used? DFS? PFS? Ext4? XFS?
 - Recommended file system filesystem **Block size**? **File system metadata**?
 - Recommended **File size on disk**?
 - **CPU selection** - number of threads , cache level, hyper threading. Cilicon
 - **Ram Selection** - and placement on chassi
 - **Network** - the difference b/w 40Gbit and 25Gbit. NIC Bonding is good for?
 - PCI Express 3, 16 Langes. PCI Switch.



Maintenance challenges: OPS Perspective

How to Optimize your hardware selection ? (Analytics on your testing data)

a. Compute intensive

- i. GPU CORE intensive
- ii. CPUcores intensive
 - 1. Frequency
 - 2. Amount of thread for concurrency

b. IO intensive?

- i. SSD?
- ii. NearLine sata?
- iii. Partitions?

c. Hardware uniformity - use same hardware all over.

d. Get rid of weak links



Maintenance challenges: DevOps Perspective

- Lightweight - python code
 - **Advantage** → focus on **generic skill** set (python coding)
 - **Disadvantage** → **reinventing the wheel**
- Fault tolerance - Scale out topology
 - **Cheap desktops** → quick recovery from Image when hardware crushes
 - **Advantage** - >
 - Cheap, low risk, pay as you go.
 - Gets 90 % of the job **DONE!**
 - **Disadvantage**
 - footprint is hell.
 - Management of desktops - requires creativity
- Rely heavily on automation feature: **reproducibility from scratch**
- **Validation automation** on infrastructure & Deployment mechanism.



Maintenance challenges: DevOps Perspective

- Infrastructure
 - **Continuous integration:** continuous build, continuous packaging, installer.
 - **Continuous Deployment:** pre flight check, remote machine (on site, private cloud, cloud)
 - **Continuous testing:** Sanity, Hourly, Daily, weekly
- Reporting and Monitoring:
 - **Extensive report server** and analytics on current testing.
 - Monitoring: **Green/Red** and system real time metrics.




Maintenance challenges: Innovation Perspective

1. Innovation on Data generation (using **GPU to generate data**)
2. Innovation on Competitor DB running time (**hashing**)
3. Innovation on Workload (one **dispatcher**, many compute nodes)
4. Innovation on saving testing data /compute time
 - a. **Hashing**
 - b. **Cacheing**
 - c. **Smart Data Creating** - no need to create 100TB CSV to generate 100TB test.
 - d. **Logs: Test results were managed on The DBMS itself :)**
5. File system **ZFS Cluster** (utilizing unused disk space, redundancy, deduplication)
6. Nvidia TK1 /**GRID GPU** → **Footprint!!!**
7. **"Virtualizing GPU"** → Footprint!!!
 - a. **Dockers**
 - b. **rCUDA**
8. **Amazon p2 gpu instances** did not exist at that time (even so still very expensive)



Building the team Challenge

- Engineering **skills** required
 - Hardware: Servers
 - IO (storage, disks, RAID, IO patten)
 - FileSystems
 - Linux , CLI, CMD, installing, compiling, GIT
 - Basic network understanding
 - High Availability, Redundancy
 - HPC
- **Coding:** python.
- **QA Big** Data mindset
- **DevOps** mindset
- **Automation** Mindset
- Systematic **Innovation** methodologies



Lesson learned (insights)

- Very hard to guarantee **100% QA coverage**
- Quality in a reasonable **cost is a huge challenge**
- Very **easy to miss milestones with QA of big data**,
- each mistake create a **huge ripple effect in timelines**
- Main **costs**:
 - **Employees** : training them.
 - **Running Time**: Coding of innovative algorithms inside testing framework,
 - **Maintenance**: Automation, Monitoring, Analyzing test results, Environments setup
 - **Hardware**: innovative DevOps , innovative OPS, research
- Most of our innovation was spent on:
 - **Doing more tests with less resources**
 - agile **improvement** of backbone
 - **Avoiding big data complexity problems in our testing!**
- **Industry tools** ? great, but not always enough. Know their philosophy...
- Sometimes you need to **get your hands dirty**
- Keep a fine **balance between business and technology**



Take away message:

- Testing simplifications:
 - **Ingress** only
 - **Egress** only
 - Ingress while egress
 - Testing matrix - useful on complex big data systems
- The QA automation team
 - coding/**scripting** skills is a MUST
 - **Engineering** skill is a MUST
 - **DevOps** understanding is a MUST
 - Allocating time for **innovation** is a MUST
 - Allocating time for **cost reduction** is a MUST
- The Automation infrastructure
 - **KISS**
 - **MVP**
 - **IS A PRODUCT** by itself. With its own Product manager and Architect



Special Thanks...

- Citi innovation Center
- [Asaf - Birenzvieg Big Data & Data Science - Israel meetup](#)
- PayPal Risk and Data solution Group



Stay in touch...

- [Omid Vahdaty](#) 
- +972-54-2384178



 **HALO**
ANALYTICS



 **jajah**
Telefonica